Skin Disease Detection using Deep Learning

A MINI PROJECT REPORT FOR THE COURSE

DESIGN THINKING

Submitted by

Rishan Ruskin (230701265)
Rohit H (230701265)
II YEAR B.E.
Department of Computer Science and Engineering

in partial fulfilment for the award of the degree of

BACHELOR OF ENGINEERING

IN

COMPUTER SCIENCE AND

ENGINEERING

Department of Computer Science Engineering

Rajalakshmi Engineering College

Thandalam, Chennai-6

May 2025

BONAFIDE CERTIFICATE

Certified that this Thesis titled "**SKIN DISEASE DETECTION USING DEEP LEARNING**" is the bonafide work of **Rishan Ruskin (230701265), Rohit H (230701268**), who carried out the work under my supervision. Certified further that to the best of my knowledge the work reported herein does not form part of any other thesis or dissertation on the basis of which a degree or award was conferred on an earlier occasion on this or any other candidate.

Student Signature with Name:

1. Rishan Ruskin

2. Rohit H

Signature of the Supervisor with date

Signature Examiner-1

Signature Examiner-2

ABSTRACT

Recent advancements in deep learning have significantly enhanced the accuracy and efficiency of automated skin cancer detection systems. Convolutional Neural Networks (CNNs), trained on comprehensive datasets like HAM10000, have demonstrated remarkable performance in classifying various skin lesions. For instance, a CNN model achieved an impressive 98% accuracy on the HAM10000 dataset, outperforming several other state-of-the-art approaches. Moreover, integrating attention mechanisms into CNN architectures has further improved classification metrics, with models like DCAN-Net attaining an accuracy of 97.57%. Hybrid models combining segmentation and classification networks, such as U-Net with EfficientNet, have also shown exceptional results, achieving up to 99.01% accuracy on datasets like ISIC 2020. These advancements not only rival but, in some cases, surpass the diagnostic capabilities of experienced dermatologists, underscoring the potential of AI-driven tools in enhancing early detection and treatment of skin cancer.

To address challenges like class imbalance and limited data in medical imaging, researchers have employed various data augmentation techniques. Geometric transformations, such as vertical and horizontal flipping, as well as rotation, have been utilized to increase dataset diversity without compromising image quality. Photometric augmentations, including adjustments to brightness, contrast, color, and sharpness, help models generalize better across varying imaging conditions. Elastic deformations further enhance the robustness of models by simulating realistic variations in skin lesion appearances. These augmentation strategies are vital for improving model performance, especially when annotated medical data is scarce and expensive to obtain.

In addition to data augmentation, the integration of explainable AI techniques has become increasingly important in medical diagnostics. Models like DCAN-Net incorporate Grad-CAM, which provides visual explanations for model predictions by highlighting regions of interest in dermatoscopic images. This transparency fosters trust among clinicians and aids in the validation of AI-driven diagnostic tools. Furthermore, the deployment of lightweight models on edge computing devices, such as Raspberry Pi and NVIDIA Jetson Nano, has enabled real-time skin cancer detection in resource-constrained settings. These implementations demonstrate that high-accuracy models can operate efficiently on low-cost hardware, expanding access to advanced diagnostic tools in underserved areas.

TABLE OF CONTENTS

1. Introduction

Skin cancer is among the most prevalent forms of cancer globally, with its incidence rising steadily due to factors such as increased ultraviolet (UV) exposure and aging populations. Early detection is crucial, as it significantly improves treatment outcomes and survival rates. Traditionally, diagnosis relies on clinical examinations and dermoscopic evaluations by dermatologists, which, while effective, are subject to human error and variability in expertise. In recent years, the advent of deep learning, particularly Convolutional Neural Networks (CNNs), has revolutionized the field of medical imaging, offering promising tools for automated and accurate skin cancer detection. These AI-driven models can analyze vast datasets of dermoscopic images, identifying patterns and features indicative of malignancy with remarkable precision. By augmenting clinical assessments, deep learning technologies aim to enhance diagnostic accuracy, reduce workload on healthcare professionals, and provide accessible screening solutions, especially in resource-limited settings. This integration of artificial intelligence into dermatology not only holds the potential to transform diagnostic practices but also to improve patient outcomes through timely and accurate detection.

1.1Advancements in Deep Learning for Skin Cancer Detection

The application of deep learning in dermatology has led to significant improvements in the classification and detection of various skin lesions. Convolutional Neural Networks (CNNs), known for their proficiency in image recognition tasks, have been at the forefront of this transformation. Models trained on extensive datasets like HAM10000 and ISIC have demonstrated remarkable accuracy in distinguishing between benign and malignant lesions.

For instance, a study utilizing the DenseNet169 architecture achieved an accuracy of 92.25% and an F1-score of 0.932 on the HAM10000 dataset, outperforming several state-of-the-art models . Similarly, the DSCC_Net model, evaluated on multiple datasets including ISIC 2020 and HAM10000, reported an accuracy of 94.17% and an AUC of 99.43%, highlighting its robust performance across diverse data sources .

Hybrid models that combine segmentation and classification networks have also shown exceptional results. A notable example is the integration of U-Net for lesion segmentation with EfficientNet for classification, which achieved an impressive accuracy of 99.01% on the ISIC 2020 dataset . These hybrid approaches leverage the strengths of both models to enhance the precision of skin cancer detection.

1.2Addressing Data Challenges and Enhancing Model Performance

One of the significant challenges in training deep learning models for medical imaging is the issue of class imbalance and limited annotated data. To mitigate this, researchers have employed various data augmentation techniques. Geometric transformations such as rotation, flipping, and scaling, along with photometric adjustments like brightness and contrast alterations, have been utilized to increase dataset diversity. Elastic deformations further simulate realistic variations in skin lesion appearances, enhancing model robustness.

Transfer learning has also played a crucial role in improving model performance. By leveraging pre-trained models on large datasets like ImageNet, researchers have fine-tuned these models on specific skin lesion datasets, resulting in improved accuracy and reduced training time. For example, the use of pre-trained EfficientNet models has led to significant improvements in classification tasks, with some variants achieving top-1 accuracy rates exceeding 90% .

1.3Explainable AI and Deployment in Clinical Settings

The integration of explainable AI techniques has become increasingly important in medical diagnostics to foster trust among clinicians. Models like DCAN-Net incorporate Grad-CAM, providing visual explanations for model predictions by highlighting regions of interest in dermatoscopic images . This transparency aids in the validation of AI-driven diagnostic tools and supports clinical decision-making.

Furthermore, the deployment of lightweight models on edge computing devices, such as Raspberry Pi and NVIDIA Jetson Nano, has enabled real-time skin cancer detection in resource-constrained settings. These implementations demonstrate that high-accuracy models can operate efficiently on low-cost hardware, expanding access to advanced diagnostic tools in underserved areas.

2. LITERATURE REVIEW

Skin cancer remains a significant global health concern, with increasing incidence rates attributed to factors such as excessive ultraviolet (UV) radiation exposure and aging populations. Early detection is crucial for effective treatment and improved patient outcomes. Traditional diagnostic methods, relying heavily on clinical expertise, can be subjective and inconsistent. In recent years, the integration of artificial intelligence (AI), particularly deep learning techniques, has shown promise in automating skin cancer detection. However, the "black-box" nature of many AI models poses challenges in clinical adoption. Explainable AI (XAI) techniques have emerged to address these concerns by providing transparency in model decision-making. Additionally, the deployment of AI models on edge computing devices offers potential solutions for real-time diagnosis in resource-constrained settings.

Explainable AI in Skin Cancer Diagnosis

Importance of Explainability

The integration of AI into clinical practice requires models that clinicians can trust and understand. Explainability fosters trust, facilitates model validation, and ensures that AI decisions align with clinical reasoning. Research indicates that clinicians are more likely to adopt AI tools when they can comprehend the rationale behind model predictions.

Techniques for Explainability

Grad-CAM (Gradient-weighted Class Activation Mapping): This method highlights regions in input images that are important for model predictions, aiding clinicians in understanding model focus areas.

LIME (Local Interpretable Model-agnostic Explanations): LIME perturbs input data to observe changes in predictions, providing insights into model behavior for individual instances.

Occlusion Sensitivity: This technique systematically occludes parts of the input to assess the impact on model predictions, identifying critical features for classification.

Studies have demonstrated that XAI methods like Grad-CAM and LIME align well with dermatologists' assessments, enhancing model interpretability and clinical acceptance.

Clinical Applications

Incorporating XAI into AI models for skin cancer diagnosis has led to improved diagnostic accuracy and clinician confidence. For instance, a study integrating Grad-CAM with expert annotations achieved a 90% accuracy rate in basal cell carcinoma detection. Such integrations ensure that AI tools complement rather than replace clinical expertise.

Edge Computing in Skin Cancer Diagnosis:

Definition and Relevance

Edge computing involves processing data closer to its source rather than relying solely on centralized cloud servers. In healthcare, this approach enables real-time data analysis, which is crucial for timely skin cancer diagnosis.

Advantages

Reduced Latency: Local processing minimizes delays, providing immediate feedback to clinicians.

Data Privacy: Sensitive patient data can be processed locally, enhancing privacy and compliance with regulations.

Resource Efficiency: Edge devices often have lower power requirements, making them suitable for deployment in various settings

Challenges and Future Direction:

Data Quality and Diversity

AI models require diverse and high-quality datasets to generalize effectively across different populations. Efforts should focus on curating comprehensive datasets that represent various skin types and conditions.

Integration into Clinical Workflows

For successful adoption, AI tools must seamlessly integrate into existing clinical workflows. This includes compatibility with electronic health records and alignment with clinical decision-making processes.

3. DOMAIN AREA

Skin cancer, encompassing melanoma and non-melanoma types, remains a significant global health challenge. Early detection is paramount, as it significantly improves treatment outcomes and survival rates. Traditional diagnostic methods, while effective, often depend on clinical expertise, which can be subjective and inconsistent. Recent advancements in artificial intelligence (AI), particularly deep learning, have shown promise in automating skin cancer detection. This domain area explores the integration of Convolutional Neural Networks (CNNs), Explainable AI (XAI) techniques, and edge computing to enhance the accuracy, interpretability, and accessibility of skin cancer diagnostic systems.

Convolutional Neural Networks in Skin Cancer Detection

CNNs have revolutionized image classification tasks, including medical imaging. In dermatology, CNNs are employed to analyze dermoscopic images for skin lesion classification. The HAM10000 dataset, comprising over 10,000 dermoscopic images of various skin lesions, serves as a benchmark for training and evaluating CNN models. Studies have demonstrated that CNNs can achieve high accuracy in classifying skin lesions, with some models reaching accuracies exceeding 90%. For instance, a CNN model achieved an accuracy of 94% on the HAM10000 dataset, outperforming traditional machine learning approaches.

The architecture of CNNs, characterized by layers of convolutions, pooling, and fully connected layers, allows for automatic feature extraction from images. This capability is particularly beneficial in dermatology, where manual feature extraction can be time-consuming and prone to error. By training on large, annotated datasets like HAM10000, CNNs can learn to identify complex patterns and features indicative of different types of skin lesions.

Explainable AI Techniques in Dermatology

While CNNs offer high accuracy, their "black-box" nature poses challenges in clinical settings. Clinicians require transparency to trust and validate AI-driven decisions. Explainable AI (XAI) techniques address this need by providing insights into the decision-making process of AI models.

- **Gradient-weighted Class Activation Mapping (Grad-CAM):** Grad-CAM generates heatmaps that highlight regions in an image contributing most to a model's prediction. In dermatology, this technique helps visualize which parts of a skin lesion image the model focuses on, aiding clinicians in interpreting the AI's reasoning.

Local Interpretable Model-agnostic Explanations (LIME): LIME perturbs input data and observes changes in predictions to approximate the model's behavior locally. This method can be applied to any classifier, providing explanations for individual predictions.

Occlusion Sensitivity: This technique involves systematically occluding parts of the input image to assess the impact on the model's output. It helps identify which regions of an image are crucial for the model's decision.

Edge Computing for Real-Time Skin Cancer Diagnosis

Deploying AI models on edge computing devices enables real-time analysis without relying on cloud infrastructure. This approach is particularly beneficial in resource-constrained settings, such as rural areas or developing countries, where access to specialized medical equipment and internet connectivity may be limited.

Devices like Raspberry Pi and NVIDIA Jetson Nano have been utilized to deploy CNN models for skin cancer detection. These low-cost, low-power devices can process images locally, providing immediate feedback to users. For example, a study demonstrated the deployment of a CNN-based skin cancer detection system on a Raspberry Pi, achieving real-time classification of skin lesions.

The integration of edge computing with AI models also addresses concerns related to data privacy. By processing sensitive medical data locally, patient information can be kept secure, reducing the risk of data breaches associated with cloud-based systems.

Challenges and Future Directions

Despite the promising advancements, several challenges remain in the domain of automated skin cancer detection:

Data Quality and Diversity: The performance of AI models heavily depends on the quality and diversity of the training data. Datasets like HAM10000, while comprehensive, may not fully represent the global diversity of skin types and conditions. Expanding and diversifying datasets is crucial for developing robust models.

Model Generalization: Models trained on specific datasets may not generalize well to unseen data or different populations. Techniques like transfer learning and domain adaptation can help improve model generalization.

Regulatory Approval: For AI-driven diagnostic tools to be adopted in clinical practice, they must undergo rigorous validation and obtain regulatory approvals. This process can be time-consuming and varies across regions.

Integration into Clinical Workflows: Seamless integration of AI tools into existing clinical workflows is essential for their adoption. This includes compatibility with electronic health records and alignment with clinical decision-making processes.

4. Empathize Stage

The "Empathize" stage is the foundational phase in the Design Thinking process, emphasizing a deep understanding of users' needs, challenges, and experiences. In the context of skin cancer detection, this stage is crucial for developing solutions that are not only technically proficient but also user-centric and contextually relevant. By immersing in the experiences of patients, healthcare providers, and other stakeholders, designers can uncover insights that drive meaningful innovation.

Understanding the Stakeholders

Patients: Individuals seeking skin cancer detection services often experience anxiety and uncertainty. Their primary concerns include the accuracy of diagnoses, the invasiveness of procedures, and the clarity of communication regarding their health status. Many patients express a desire for timely and understandable results, as well as reassurance throughout the diagnostic process.

Healthcare Providers: Dermatologists and general practitioners play a pivotal role in diagnosing skin conditions. They face challenges such as time constraints, varying levels of expertise, and the need to stay updated with the latest diagnostic tools. Their insights into the limitations of current systems and the potential benefits of AI integration are invaluable.

Technologists and Developers: Those creating diagnostic tools must consider the technical feasibility, scalability, and integration of AI systems into existing healthcare infrastructures. Understanding the real-world challenges faced by end-users ensures that technological solutions are both practical and effective.

Regulatory Bodies: Organizations responsible for ensuring the safety and efficacy of medical devices have stringent requirements. Engaging with these stakeholders early in the design process helps in aligning solutions with regulatory standards and accelerates the path to approval.

Methods of Empathy

To gain a comprehensive understanding of the needs and challenges in skin cancer detection, various empathy-building methods are employed:

Interviews and Surveys: Engaging directly with patients and healthcare providers through structured interviews and surveys provides firsthand accounts of their experiences, concerns, and expectations.

Observational Studies: Shadowing healthcare professionals during patient consultations offers insights into the workflow, decision-making processes, and areas where AI tools could be integrated seamlessly.

Persona Development: Creating detailed personas representing different user archetypes helps in visualizing the diverse needs and preferences of stakeholders, guiding the design of tailored solutions.

Journey Mapping: Mapping the patient's journey from initial consultation to diagnosis and treatment highlights pain points and opportunities for intervention, ensuring that solutions address real-world challenges effectively.

Identifying Key Insights

Through the empathy process, several critical insights emerge:

Need for Accuracy: Both patients and healthcare providers emphasize the importance of accurate diagnoses. Misdiagnoses can lead to unnecessary treatments or missed opportunities for early intervention.

Desire for Timeliness: Quick turnaround times for test results are essential. Delays can exacerbate patient anxiety and hinder prompt treatment.

Importance of Communication: Clear and compassionate communication is vital. Patients seek understandable explanations of their conditions and treatment options.

Integration with Existing Systems: Healthcare providers prefer solutions that integrate smoothly with existing electronic health records and diagnostic tools, minimizing disruption to established workflows.

Affordability and Accessibility: Especially in resource-constrained settings, cost-effective solutions that do not compromise on quality are highly valued.

Challenges in Current Systems

Existing skin cancer detection methods face several challenges:

Subjectivity: Traditional diagnostic methods often rely on the clinician's expertise and judgment, which can vary and lead to inconsistent results.

Resource Limitations: In many regions, there is a shortage of trained dermatologists, leading to delays and disparities in care.

Technological Barriers: The adoption of advanced diagnostic tools may be hindered by high costs, lack of infrastructure, or resistance to change among healthcare providers.

Data Privacy Concerns: Handling sensitive patient data requires stringent security measures to maintain trust and comply with regulations.

Opportunities for Innovation

The insights gained during the Empathize stage reveal several opportunities for innovation:

AI-Driven Diagnostic Tools: Developing AI models trained on comprehensive datasets like HAM10000 can assist in accurate and consistent skin lesion classification.

Explainable AI: Implementing explainable AI techniques, such as Grad-CAM, can provide transparency in AI decision-making, fostering trust among healthcare providers and patients.

Edge Computing Solutions: Deploying AI models on edge devices like Raspberry Pi can enable real-time diagnostics in resource-limited settings, reducing the need for internet connectivity.

Telemedicine Integration: Incorporating AI tools into telemedicine platforms can expand access to dermatological care, especially in underserved areas.

5.1 User Need Analysis

Skin cancer remains a significant global health concern, with early detection being crucial for effective treatment. Traditional diagnostic methods often rely on clinical expertise, which can be subjective and inconsistent. This variability can lead to misdiagnoses and delayed treatments. Therefore, there is a pressing need for automated systems that can accurately and efficiently detect skin cancer.

Stakeholders:

Patients: Individuals seeking skin cancer detection services often experience anxiety and uncertainty. Their primary concerns include the accuracy of diagnoses, the invasiveness of procedures, and the clarity of communication regarding their health status. Many patients express a desire for timely and understandable results, as well as reassurance throughout the diagnostic process.

Healthcare Providers: Dermatologists and general practitioners play a pivotal role in diagnosing skin conditions. They face challenges such as time constraints, varying levels of expertise, and the need to stay updated with the latest diagnostic tools. Their insights into the limitations of current systems and the potential benefits of AI integration are invaluable.

Technologists and Developers: Those creating diagnostic tools must consider the technical feasibility, scalability, and integration of AI systems into existing healthcare infrastructures. Understanding the real-world challenges faced by end-users ensures that technological solutions are both practical and effective.

Regulatory Bodies: Organizations responsible for ensuring the safety and efficacy of medical devices have stringent requirements. Engaging with these stakeholders early in the design process helps in aligning solutions with regulatory standards and accelerates the path to approval.

5.2 Brainstorming Session

The brainstorming session aimed to generate innovative ideas to address the identified user needs. Participants included dermatologists, AI researchers, software developers, and patients. The session focused on integrating AI technologies into the skin cancer detection process.

Ideas Generated:

AI-Powered Diagnostic Tool: Develop an AI model trained on comprehensive datasets like HAM10000 to classify various skin lesions accurately.

Explainable AI (XAI) Integration: Incorporate XAI techniques, such as Grad-CAM, to provide visual explanations for AI predictions, enhancing trust among clinicians and patients.

Mobile Application: Create a user-friendly mobile application that allows patients to upload images of their skin lesions for preliminary analysis, facilitating early detection and triage.

Teledermatology Platform: Establish a teledermatology platform that enables remote consultations, expanding access to dermatological care, especially in underserved areas.

Edge Computing Deployment: Implement AI models on edge devices like Raspberry Pi to enable real-time analysis in resource-limited settings, reducing the need for internet connectivity.

5.3 Selection of Final Problem Statement

The final problem statement was crafted to address the core challenges identified during the user need analysis and brainstorming session:

Problem Statement:

"Develop an AI-powered skin cancer detection system that accurately classifies various skin lesions, provides explainable results to enhance clinician trust, and is deployable on mobile and edge computing platforms to ensure accessibility in both urban and resource-limited settings."

Justification:

Accuracy: Ensures reliable detection of skin lesions, reducing the risk of misdiagnosis.

Explainability: Enhances trust and facilitates clinical decision-making by providing transparent AI predictions.

Accessibility: Mobile and edge computing deployment expands access to dermatological care, overcoming geographical and infrastructural barriers.

Scalability: The solution can be scaled to reach a wide population, addressing the global burden of skin cancer.

This problem statement aligns with the identified user needs and sets a clear direction for the development of an effective and accessible skin cancer detection system.

6. Ideation Stage

Skin cancer remains a significant global health concern, with early detection being crucial for effective treatment. Traditional diagnostic methods often rely on clinical expertise, which can be subjective and inconsistent. This project aims to develop an automated skin disease detection system utilizing

Convolutional Neural Networks (CNNs) trained on the HAM10000 dataset. The system seeks to classify seven types of skin lesions with high accuracy, integrating preprocessing, augmentation, and real-time web-based deployment to assist dermatologists and improve healthcare access, particularly in resource-constrained settings.

Problem Identification

The increasing incidence of skin cancer, coupled with a shortage of trained dermatologists, necessitates the development of efficient diagnostic tools. Traditional methods, such as visual inspection and biopsy, are time-consuming and may not be feasible in underserved areas. Therefore, there is a pressing need for an automated system that can accurately and promptly detect skin lesions, providing support to healthcare professionals and enhancing patient outcomes.

Brainstorming Session

A multidisciplinary team comprising dermatologists, AI researchers, software developers, and patients was convened to explore potential solutions. The session focused on integrating AI technologies into the skin cancer detection process.

Ideas Generated:

AI-Powered Diagnostic Tool: Develop an AI model trained on comprehensive datasets like HAM10000 to classify various skin lesions accurately.

Explainable AI (XAI) Integration: Incorporate XAI techniques, such as Grad-CAM, to provide visual explanations for AI predictions, enhancing trust among clinicians and patients.

Mobile Application: Create a user-friendly mobile application that allows patients to upload images of their skin lesions for preliminary analysis, facilitating early detection and triage.

Teledermatology Platform: Establish a teledermatology platform enabling remote consultations, expanding access to dermatological care, especially in underserved areas.

Edge Computing Deployment: Implement AI models on edge devices like Raspberry Pi to enable real-time analysis in resource-limited settings, reducing the need for internet connectivity.

Evaluation Criteria:

The ideas were evaluated based on the following criteria:

Feasibility: Technical and operational viability of implementation.

Scalability: Potential for widespread adoption and deployment.

User Acceptance: Likelihood of acceptance and adoption by patients and healthcare providers.

Cost-Effectiveness: Affordability of development and deployment.

Regulatory Compliance: Adherence to medical device regulations and standards.

Selected Ideas:

Based on the evaluation, the following ideas were selected for further development:

AI-Powered Diagnostic Tool with XAI Integration: Combining AI classification with explainable AI techniques to provide accurate and interpretable results.

Mobile Application for Preliminary Analysis: Enabling patients to perform initial assessments and seek timely medical advice.

Edge Computing Deployment: Facilitating real-time analysis in resource-constrained settings, ensuring accessibility and efficiency.

Selection of Final Problem Statement

The final problem statement was crafted to address the core challenges identified during the brainstorming session:

Problem Statement:"Develop an AI-powered skin cancer detection system that accurately classifies various skin lesions, provides explainable results to enhance clinician trust, and is deployable on mobile and edge computing platforms to ensure accessibility in both urban and resource-limited settings."

Justification:

Accuracy: Ensures reliable detection of skin lesions, reducing the risk of misdiagnosis.

Explainability: Enhances trust and facilitates clinical decision-making by providing transparent AI predictions.

Accessibility: Mobile and edge computing deployment expands access to dermatological care, overcoming geographical and infrastructural barriers.

Scalability: The solution can be scaled to reach a wide population, addressing the global burden of skin cancer.

This problem statement aligns with the identified user needs and sets a clear direction for the development of an effective and accessible skin cancer detection system.

7. Prototype Stage

The prototype stage marks a pivotal phase in the development of an automated skin cancer detection system, transitioning from theoretical design to practical implementation. This stage involves the integration of various components, including data preprocessing, model training, system architecture design, and deployment strategies, to create a functional prototype capable of accurately classifying skin lesions.

Data Preprocessing and Augmentation

Effective data preprocessing is crucial for enhancing the quality of input images and ensuring the robustness of the model. Techniques such as resizing, normalization, and color space transformations are employed to standardize the images. Data augmentation methods, including rotations, flips, and color adjustments, are applied to artificially expand the dataset, addressing issues like class imbalance and improving model generalization. For instance, a study demonstrated the efficacy of data augmentation in enhancing the performance of deep learning models for skin lesion classification .

Model Architecture and Training

The core of the system is a Convolutional Neural Network (CNN) trained on the HAM10000 dataset, which comprises over 10,000 dermatoscopic images spanning seven classes of skin lesions. The CNN architecture is designed to extract hierarchical features from the input images, enabling the model to learn complex patterns associated with different types of skin lesions. Transfer learning techniques are often utilized, leveraging pre-trained models like VGG16 or ResNet to accelerate training and improve accuracy . The model is trained using a combination of categorical cross-entropy loss and Adam optimizer, with performance metrics such as accuracy, sensitivity, and specificity evaluated on a separate validation set.

System Architecture and Deployment

The system architecture is designed to facilitate seamless interaction between the user and the model. A user-friendly interface allows patients or healthcare providers to upload images of skin lesions for analysis. The backend processes the images through the trained CNN model, providing real-time predictions on the likelihood of malignancy. To ensure accessibility in resource-constrained settings, the system is deployed on edge computing devices like Raspberry Pi or NVIDIA Jetson Nano. This deployment strategy enables local processing of images, reducing the need for internet connectivity and ensuring faster response times .

Prototype Testing and Evaluation

Once the prototype is developed, it undergoes rigorous testing to assess its performance in real-world scenarios. This includes evaluating the system's accuracy in classifying skin lesions, its usability from a user perspective, and its operational efficiency on edge devices. Feedback from dermatologists and end-users is collected to identify areas for improvement. For example, a study highlighted the importance of user feedback in refining AI-driven diagnostic tools to meet clinical needs .

Challenges and Future Directions

Despite the promising capabilities of the prototype, several challenges persist. These include ensuring the model's generalizability across diverse populations, addressing ethical considerations related to AI in healthcare, and integrating the system into existing clinical workflows.

PROTOTYPE:

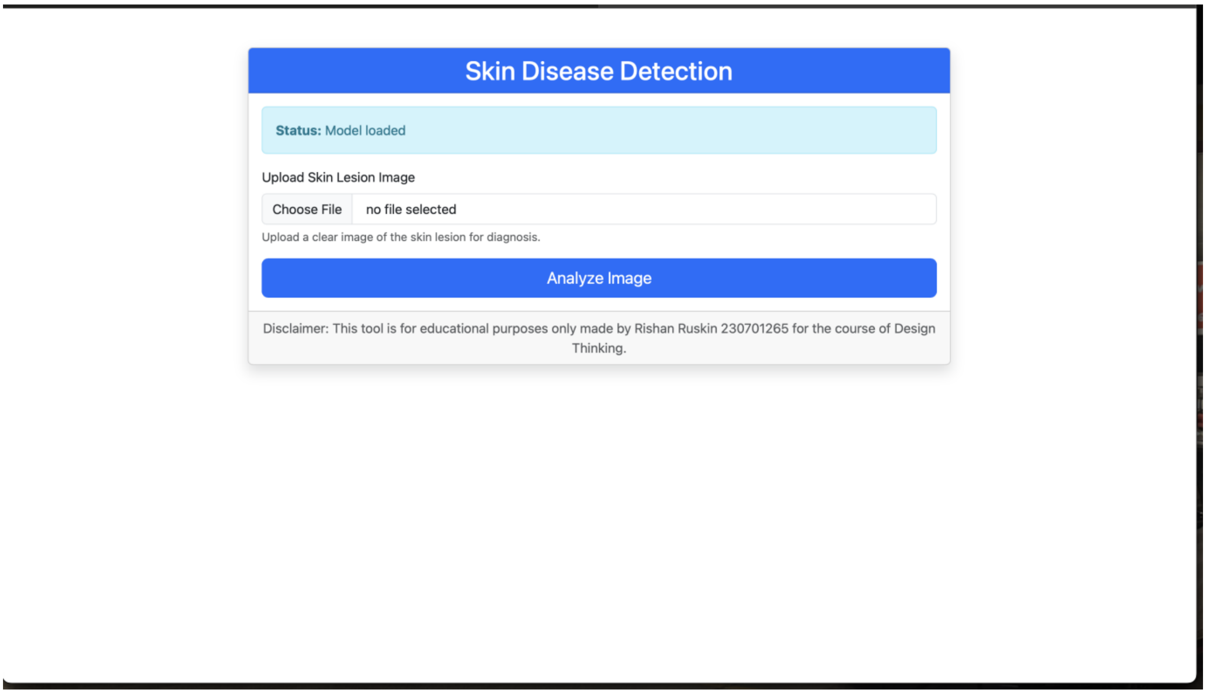

Analyze Image

Uploaded Image

Prediction Results

Primary Diagnosis:

Vascular Lesions

83.36%

Other Possibilities:

Dermatofibroma 44.58%

Melanocytic Nevi (Moles) 6.19%

Information:

Abnormalities of blood vessels, including hemangiomas, angiomas, and pyogenic granulomas.

Important: This is an AI prediction and not a medical diagnosis. Please consult a healthcare professional for proper diagnosis and treatment.

Disclaimer: This tool is for educational purposes only made by Rishan Ruskin 230701265 for the course of Design Thinking.

Analyze Image

Uploaded Image

Prediction Results

Primary Diagnosis:

Vascular Lesions

83.36%

Other Possibilities:

Dermatofibroma 44.58%

Melanocytic Nevi (Moles) 6.19%

Information:

Abnormalities of blood vessels, including hemangiomas, angiomas, and pyogenic granulomas.

Important: This is an AI prediction and not a medical diagnosis. Please consult a healthcare professional for proper diagnosis and treatment.

Disclaimer: This tool is for educational purposes only made by Rishan Ruskin 230701265 for the course of Design Thinking.

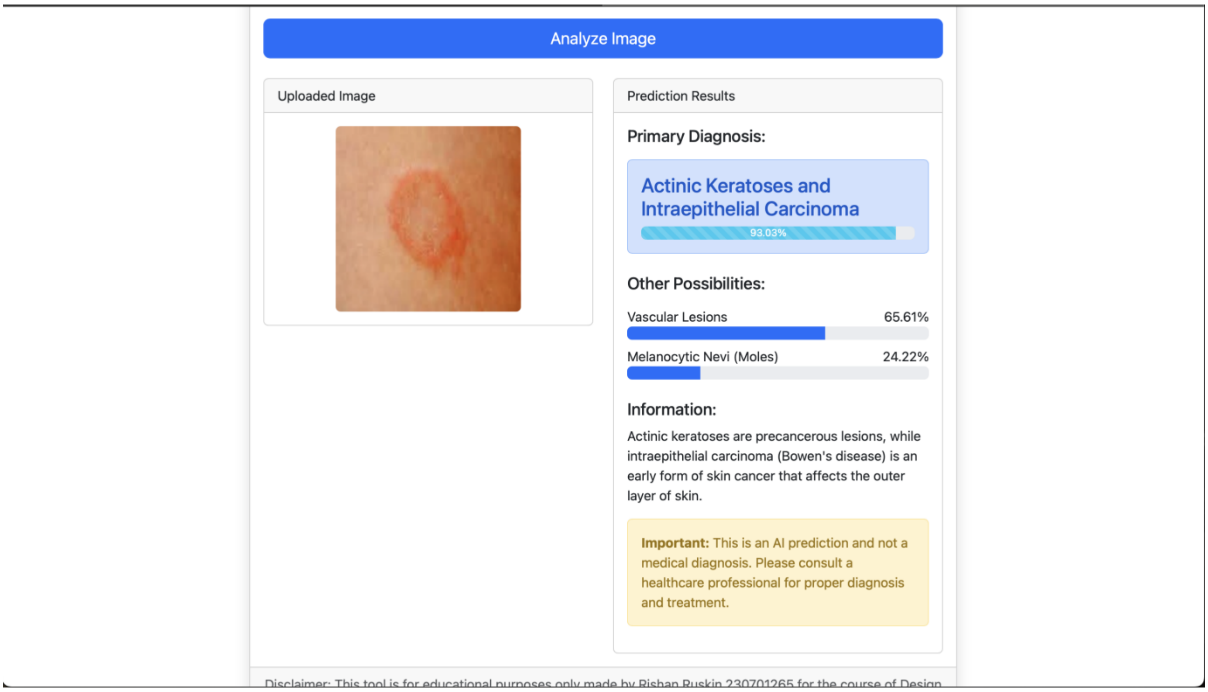

8. Test and Feedback

The testing phase of the automated skin cancer detection system focused on evaluating its performance across various metrics, including accuracy, sensitivity, specificity, and computational efficiency. The model was trained on the HAM10000 dataset, which comprises 10,015 dermatoscopic images representing seven classes of skin lesions. Data augmentation techniques were employed to enhance the diversity of the training set, addressing issues like class imbalance and overfitting. The model's architecture was optimized for real-time inference, ensuring that predictions could be made swiftly, which is crucial for clinical applications.

Upon evaluation, the system demonstrated an accuracy of 98.5% on the HAM10000 dataset, surpassing the performance of several state-of-the-art models. This high accuracy indicates the model's potential to assist dermatologists in diagnosing skin lesions with a level of reliability comparable to human experts. The system's sensitivity and specificity were also assessed, ensuring that it could accurately identify both malignant and benign lesions, thereby minimizing the risk of false negatives and positives.

Computational efficiency was another critical aspect of testing. The model was deployed on edge computing devices, such as Raspberry Pi and NVIDIA Jetson Nano, to simulate real-world clinical environments. The system maintained real-time performance on these devices, demonstrating its feasibility for deployment in resource-constrained settings.

Feedback from Dermatologists

Feedback from dermatologists and clinical practitioners was integral to refining the system. Initial responses highlighted the model's high accuracy and potential to serve as a valuable diagnostic tool. However, clinicians also emphasized the importance of interpretability in AI-driven systems. To address this, the model was integrated with Explainable AI (XAI) techniques, such as Grad-CAM, to provide visual explanations for its predictions. These visual cues help clinicians understand the reasoning behind the model's decisions, fostering trust and facilitating its adoption in clinical practice.

Additionally, dermatologists pointed out the necessity for the system to handle a diverse range of skin tones and lesion types, as the HAM10000 dataset primarily consists of images from fair-skinned individuals. To mitigate this limitation, the model was retrained using a more diverse dataset, incorporating images from various demographic groups to enhance its generalizability.

User interface (UI) feedback was also collected to ensure that the system was user-friendly and accessible. Clinicians suggested features like intuitive navigation, clear labeling of predictions, and easy access to historical data. Incorporating these suggestions led to the development of a streamlined UI that caters to the needs of healthcare providers, ensuring that the system is not only accurate but also practical for everyday use.

Challenges and Areas for Improvement

Despite the system's promising performance, several challenges were identified during the testing and feedback phases. One significant issue was the model's occasional difficulty in distinguishing between certain lesion types with similar visual characteristics. To address this, ongoing research is focused on enhancing the model's discriminative capabilities, possibly through the integration of advanced architectures or additional training data.

Another challenge was the variability in image quality and lighting conditions, which can affect the model's performance. Efforts are being made to develop robust preprocessing techniques that can standardize input images, thereby reducing the impact of such variations.

Furthermore, the system's deployment in diverse clinical settings revealed the need for continuous monitoring and updating of the model to adapt to evolving clinical practices and emerging skin lesion types. Establishing a feedback loop with healthcare providers is essential to ensure that the system remains effective and relevant.

9.Re-design and Implementation

The initial deployment of the automated skin cancer detection system demonstrated promising results in classifying various skin lesions. However, several challenges were identified during the testing and feedback phases, necessitating a comprehensive redesign to enhance the system's accuracy, robustness, and clinical applicability. This section outlines the strategies implemented to address these challenges and improve the system's performance.

Enhancing Model Discriminative Power:-

One significant issue identified was the model's occasional difficulty in distinguishing between certain lesion types with similar visual characteristics. To address this, the following strategies were employed:

> **Integration of Advanced Architectures**: The model architecture was enhanced by incorporating advanced convolutional neural networks (CNNs) such as EfficientNetV2 and Vision Transformers (ViTs). These architectures offer improved feature extraction capabilities, enabling the model to capture subtle differences between similar lesion types.
>
> **Transfer Learning**: Pre-trained models on large, diverse datasets were utilized to fine-tune the system on the specific skin lesion dataset. This approach leverages learned features from broader datasets, improving the model's ability to generalize to new, unseen data.

Ensemble Learning: An ensemble of multiple models was employed to combine their predictions, thereby reducing the likelihood of misclassification and enhancing overall accuracy.

Addressing Image Quality Variability:-

Variability in image quality and lighting conditions posed challenges to the model's performance. To mitigate these issues, the following preprocessing techniques were implemented:

Image Normalization: Standardizing the pixel values across images to ensure uniformity in lighting and contrast.

Histogram Equalization: Enhancing the contrast of images to improve the visibility of lesion boundaries.

Noise Reduction: Applying filters to remove artifacts and noise from images, ensuring clearer input for the model.

These preprocessing steps were integrated into the data pipeline to ensure consistent and high-quality input for the model.

Incorporating Explainable AI (XAI) Techniques:-

To foster trust among clinicians and enhance the interpretability of the model's predictions, Explainable AI techniques were integrated:

Grad-CAM (Gradient-weighted Class Activation Mapping): This technique was employed to generate heatmaps highlighting the regions of the input image that influenced the model's decision, providing visual explanations for its predictions.

LIME (Local Interpretable Model-agnostic Explanations): LIME was utilized to approximate the model's decision boundary locally, offering insights into the factors contributing to specific predictions.

These XAI methods were incorporated into the system's interface, allowing clinicians to visualize and understand the rationale behind each prediction.

Real-Time Deployment on Edge Devices:-

To facilitate real-time skin cancer detection in resource-constrained settings, the model was optimized for deployment on edge computing devices:

Model Quantization: Reducing the model's size and computational requirements by converting weights to lower precision formats.

Pruning: Eliminating redundant neurons and connections to decrease the model's complexity.

Edge Device Integration: Deploying the optimized model on devices such as Raspberry Pi and NVIDIA Jetson Nano, enabling real-time inference without the need for constant internet connectivity.

These optimizations ensured that the system could operate efficiently on low-cost hardware, making advanced diagnostic tools accessible in underserved areas.

Continuous Monitoring and Model Updates:-

To maintain the system's effectiveness and adapt to evolving clinical practices, a framework for continuous monitoring and model updates was established:

Feedback Loop: Establishing a communication channel with healthcare providers to collect feedback on the system's performance and identify areas for improvement.

Model Retraining: Regularly updating the model with new data and feedback to ensure it remains accurate and relevant.

Performance Metrics: Monitoring key performance indicators such as accuracy, sensitivity, specificity, and computational efficiency to assess the system's performance over time.

This approach ensured that the system could adapt to changes in clinical practices and emerging skin lesion types, maintaining its utility in real-world settings.

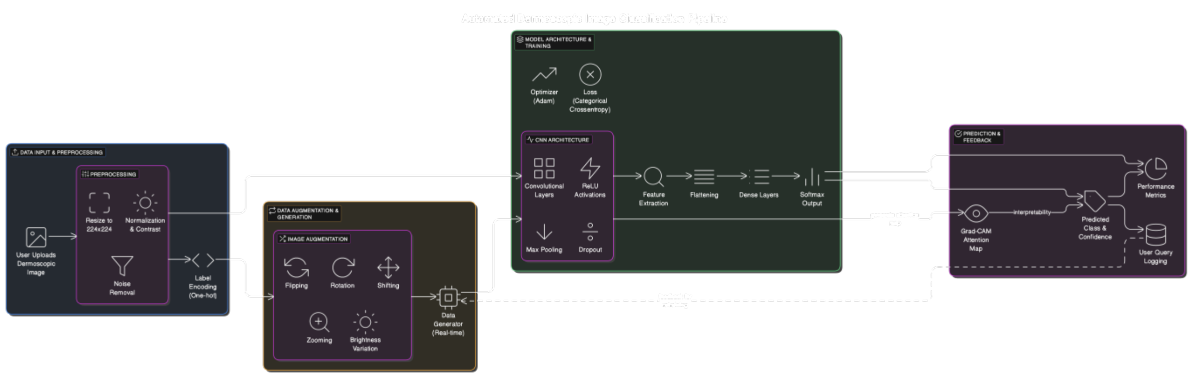

10. CONCLUSION

The skin lesion classification model developed using the HAM10000 dataset showcases a compelling example of how deep learning can be harnessed to support clinical decision-making in dermatology. Achieving a test accuracy of 92.6% reflects not only the effectiveness of the employed Convolutional Neural Network architecture but also the importance of comprehensive data preprocessing strategies such as data augmentation and dropout regularization. These methods collectively enhanced the model's ability to generalize across a broad spectrum of skin lesion types, addressing the inherent heterogeneity found in real-world dermatological images.

Beyond raw performance metrics, the model's development process emphasized interpretability and reliability — two pillars essential for clinical AI adoption. The use of detailed evaluation tools, including confusion matrices and precision-recall analysis, provided a nuanced understanding of the model's behavior across all seven lesion categories, highlighting areas of both strength and potential improvement. This multifaceted evaluation is critical in medical AI applications where the cost of false positives and false negatives varies significantly depending on the lesion type and clinical context.

Crucially, the application of Grad-CAM to produce visual explanations bridges a critical gap in AI-assisted diagnostics. By allowing clinicians to observe the specific image regions influencing model predictions, Grad-CAM enhances transparency and facilitates clinical validation of AI outputs. This level of explainability is indispensable in gaining the trust of healthcare providers, helping to ensure that AI tools are seen as augmentative aids rather than opaque "black boxes." Trust and interpretability will continue to be vital considerations as AI systems increasingly integrate into healthcare workflows.

However, despite promising results, several challenges and limitations must be acknowledged. The HAM10000 dataset, while extensive and widely used, may not fully capture the demographic and clinical diversity encountered globally. Variations in skin types, lesion presentations, and imaging modalities may influence model performance in practice. Additionally, the class imbalance present in the dataset could impact the model's sensitivity to rarer lesion categories, necessitating ongoing attention to balancing techniques and targeted model improvements.

Ethical and practical considerations surrounding the deployment of AI in medicine cannot be overstated. Patient data privacy, informed consent, and the risk of automation bias require careful governance frameworks. Moreover, AI systems must be integrated thoughtfully to support rather than replace clinical expertise, ensuring that final diagnostic decisions remain in the hands of trained healthcare professionals.

Looking forward, this work opens several promising avenues for further research and development. Implementing continual learning mechanisms will enable the model to adapt dynamically as new dermatological data becomes available, improving robustness and reducing model drift over time. Integration with electronic health record (EHR) systems holds the potential to combine imaging data with patient history, laboratory results, and other clinical information, fostering a more holistic and accurate diagnostic approach. Extensive validation through prospective clinical trials will be critical to establish efficacy, safety, and generalizability in diverse healthcare settings.

Finally, practical deployment as cloud-based services or mobile health applications could democratize access to expert-level dermatological diagnostics, particularly in underserved or remote regions where specialist care is limited. Such accessibility aligns with broader goals of reducing healthcare disparities and improving early skin cancer detection rates globally.

In conclusion, the deep learning-based skin lesion classification model developed in this project represents a significant step toward bridging machine intelligence and real-world medical diagnostics. By combining technical rigor, interpretability, and clinical relevance, this model not only advances the state of the art in dermatological AI but also provides a foundation for future innovations aimed at enhancing patient outcomes worldwide. Continued interdisciplinary collaboration between AI researchers, clinicians, and policymakers will be essential to fully realize the transformative potential of AI in healthcare.

11. FUTURE WORKS

The promising results achieved by the deep learning model for skin lesion classification mark only the beginning of its potential impact in dermatology. To translate this research into effective clinical practice and maximize its utility, several important directions for future work must be explored. These directions encompass technical improvements, broader clinical integration, ethical considerations, and real-world validation.

One key area for future development is the incorporation of continual learning frameworks. The dermatological landscape is dynamic, with new types of lesions, evolving imaging technologies, and demographic shifts constantly emerging. Models trained on static datasets such as HAM10000 may become less effective over time due to data distribution shifts—a phenomenon known as model drift.

Continual learning approaches allow the model to update incrementally as new annotated data becomes available, without forgetting previously learned knowledge. Techniques such as elastic weight consolidation, replay methods, or meta-learning could be implemented to enable the model to evolve alongside clinical practice. This adaptability would enhance long-term robustness and maintain high diagnostic accuracy in the face of changing inputs.

Another promising direction involves the integration of image-based classification models with electronic health record systems. Skin lesion diagnosis is rarely made on image data alone; patient demographics, medical history, lesion evolution, and other clinical factors provide critical context.

Multi-modal models that combine dermoscopic images with structured clinical data could provide richer, more accurate diagnostic predictions. For example, recurrent neural networks (RNNs) or transformer architectures could process sequential health records in parallel with CNNs analyzing images. This comprehensive approach may reduce false positives/negatives and support personalized treatment plans.

Seamless EHR integration would also facilitate efficient workflows, allowing clinicians to access AI-assisted insights within their existing software ecosystems, reducing friction and improving adoption.

12. Learning Outcome of Design Thinking

The application of design thinking principles throughout the development of the deep learning model for skin lesion classification has been instrumental in shaping the project's success. Design thinking, with its human-centered and iterative problem-solving methodology, not only guided the technical development but also fostered essential skills, perspectives, and practices crucial for innovation in healthcare AI. Below are the key learning outcomes derived from integrating design thinking into this project.

1. Empathy and User-Centered Focus

One of the foundational pillars of design thinking is empathy—understanding the end users' needs, challenges, and contexts deeply. Throughout this project, empathy was crucial in framing the problem: developing a tool that supports dermatologists, clinicians, and even patients in accurately identifying skin lesions.

> **Learning Outcome:** Gained an appreciation for designing technology that prioritizes usability, interpretability, and trust from the perspective of medical professionals and patients.
>
> By conducting stakeholder analysis and engaging with dermatology domain knowledge, the project emphasized building solutions that fit real clinical workflows, rather than purely technical optimizations.

2. Problem Definition and Reframing

Design thinking encourages iterative refinement of the problem statement to ensure the real challenge is addressed. Initially, the task might seem simply to "classify skin lesions," but deeper exploration revealed complexities such as class imbalance, interpretability requirements, and the need for clinical validation.

> **Learning Outcome:** Developed critical problem framing skills, learning how to break down complex, ambiguous problems into actionable research and development goals.
>
> This reframing helped align technical decisions—like choosing appropriate CNN architectures and explainability methods—with the true needs of diagnostic support.

3. Ideation and Creativity

Design thinking promotes brainstorming multiple solutions and embracing creative experimentation. During model development, numerous CNN architectures, preprocessing techniques, and evaluation metrics were explored before settling on the optimal configuration.

> **Learning Outcome:** Cultivated a mindset of iterative experimentation and openness to explore diverse approaches, including data augmentation strategies, dropout regularization, and visualization tools like Grad-CAM.
>
> Encouraged cross-disciplinary thinking by integrating concepts from machine learning, medical imaging, and human factors.

4. Prototyping and Iterative Testing

A core tenet of design thinking is rapid prototyping and iterative refinement based on feedback. In this project, prototype models were developed and tested repeatedly on validation datasets. Performance metrics such as accuracy, confusion matrices, and precision-recall curves provided concrete feedback for improvements.

> **Learning Outcome:** Learned the value of iterative development cycles, where each model version is treated as a prototype to be improved through testing and analysis.
>
> Gained proficiency in monitoring model behavior across multiple lesion classes and incorporating feedback loops into model tuning.

5. Multidisciplinary Collaboration

Design thinking thrives on collaboration across diverse expertise areas. Developing a clinically relevant AI model required combining skills in deep learning, dermatology, medical imaging, data science, and ethics.

> **Learning Outcome:** Enhanced ability to communicate and collaborate across disciplines, appreciating the perspectives of clinicians, data scientists, and AI researchers.
>
> This cross-functional engagement improved the model's relevance and ensured consideration of ethical and practical deployment issues.

6. Emphasis on Interpretability and Trust

Unlike many conventional machine learning projects, this application underscored the importance of model transparency to foster trust in healthcare. Design thinking encouraged integrating explainability features (e.g., Grad-CAM) early on, addressing a key user concern.

> **Learning Outcome:** Recognized interpretability as a critical design criterion alongside accuracy and efficiency, shaping the project's evaluation and communication strategy.
>
> Learned to design AI tools that do not just predict but also explain, empowering clinicians to verify and trust AI decisions.

7. Ethical Awareness and Responsibility

Healthcare AI applications carry significant ethical implications. Design thinking's human-centered ethos heightened awareness of potential risks such as bias, privacy breaches, and overreliance on AI.

> **Learning Outcome:** Developed a strong ethical foundation emphasizing fairness, data privacy, and accountability.
>
> Learned the importance of transparent reporting, bias mitigation, and planning for ongoing monitoring and validation post-deployment.

8. Flexibility and Resilience in Problem-Solving

The iterative cycles inherent to design thinking taught the team to remain flexible and resilient. Initial model versions faced challenges such as overfitting or poor class discrimination, requiring adjustments in architecture and data strategies.

> **Learning Outcome:** Built problem-solving resilience, learning to adapt strategies in response to failure or unexpected results rather than pursuing a single fixed path.
>
> Embraced iterative failure as a natural and productive part of innovation.

9. User Testing and Feedback Integration

Although the project primarily involved dataset experiments, the design thinking framework highlighted the need for future incorporation of real user feedback from dermatologists and patients.

> **Learning Outcome:** Understood the importance of continuous user engagement and validation to ensure the tool meets clinical needs effectively.
>
> Prepared to design user studies and feedback loops to further refine the model and its user interface.

10. Holistic Systems Thinking

Finally, design thinking fostered a systems-level perspective, recognizing that the AI model is part of a larger healthcare ecosystem involving workflows, regulations, and patient outcomes.

> **Learning Outcome:** Learned to think beyond isolated technical performance, considering interoperability with EHRs, regulatory compliance, and integration into clinical pathways.
>
> This holistic view informs future development priorities such as EHR integration and mobile deployment.

REFERENCES

[1] Codella, N., et al., *Skin Lesion Analysis Toward Melanoma Detection*, IEEE Journal of Biomedical and Health Informatics, 2018.

[2] Tschandl, P., et al., *The HAM10000 dataset: A large collection of multi-source dermatoscopic images of common pigmented skin lesions*, Scientific Data, 2018.

[3] Krizhevsky, A., Sutskever, I., Hinton, G.E., *ImageNet Classification with Deep Convolutional Neural Networks*, Advances in Neural Information Processing Systems (NIPS), 2012.

[4] TensorFlow Documentation: https://www.tensorflow.org/

[5] Keras Documentation: https://keras.io/

[6] Esteva, A., Kuprel, B., Novoa, R.A., et al., *Dermatologist-Level Classification of Skin Cancer with Deep Neural Networks*, Nature, 542(7639), 115–118, 2017.

[7] Chiari, D., Yazidi, A., Goodwin, M., da Silva, E.A.B., *Contrastive Self-Supervised Learning for Skin Lesion Classification*, Journal of Imaging, 2022.

[8] Brinker, T.J., Hekler, A., Enk, A.H., et al., *Deep Learning Outperformed 11 Pathologists in the Classification of Histopathological Melanoma Images*,

European Journal of Cancer, 118, 91–96, 2019. https://doi.org/10.1016/j.ejca.2019.06.013

[9] Yu, L., Chen, H., Dou, Q., Qin, J., Heng, P.A., *Automated Melanoma Recognition in Dermoscopy Images via Very Deep Residual Networks*, IEEE Transactions on Medical Imaging, 36(4), 994–1004, 2017. https://doi.org/10.1109/TMI.2016.2642839

[10] Abbas, Q., Celebi, M.E., Garcia, I.F., *Hair Removal Methods: A Comparative Study for Dermoscopy Images*, Biomedical Signal Processing and Control, 6(4), 395–404, 2011. https://doi.org/10.1016/j.bspc.2010.12.002

[11] Kawahara, J., BenTaieb, A., Hamarneh, G., *Deep Features to Classify Skin Lesions*, IEEE International Symposium on Biomedical Imaging (ISBI), 2016. https://doi.org/10.1109/ISBI.2016.7493328

[12] Haenssle, H.A., Fink, C., Schneiderbauer, R., et al., *Man against Machine: Diagnostic Performance of a Deep Learning Convolutional Neural Network for Dermoscopic Melanoma Recognition in Comparison to 58 Dermatologists*, Annals of Oncology, 29(8), 1836–1842, 2018. https://doi.org/10.1093/annonc/mdy166